


@WINTER SCHOOL



@CSE-IIT Kanpur*

22nd December 2020

*Joining IIT Bombay in May 2021



MICROARCHITECTURE-SECURITY@WINTER SCHOOL



@CSE-IIT Kanpur*

22nd December 2020

*Joining IIT Bombay in May 2021

Disclaimer



Talk is not about **leaking information from your machines through MS teams.**

I am not trying for the same. Trust me 😊


After the talk, if your password is compromised, It is not me 😊

Let's start the talk then ...



Architecture:101



Microarchitecture:101

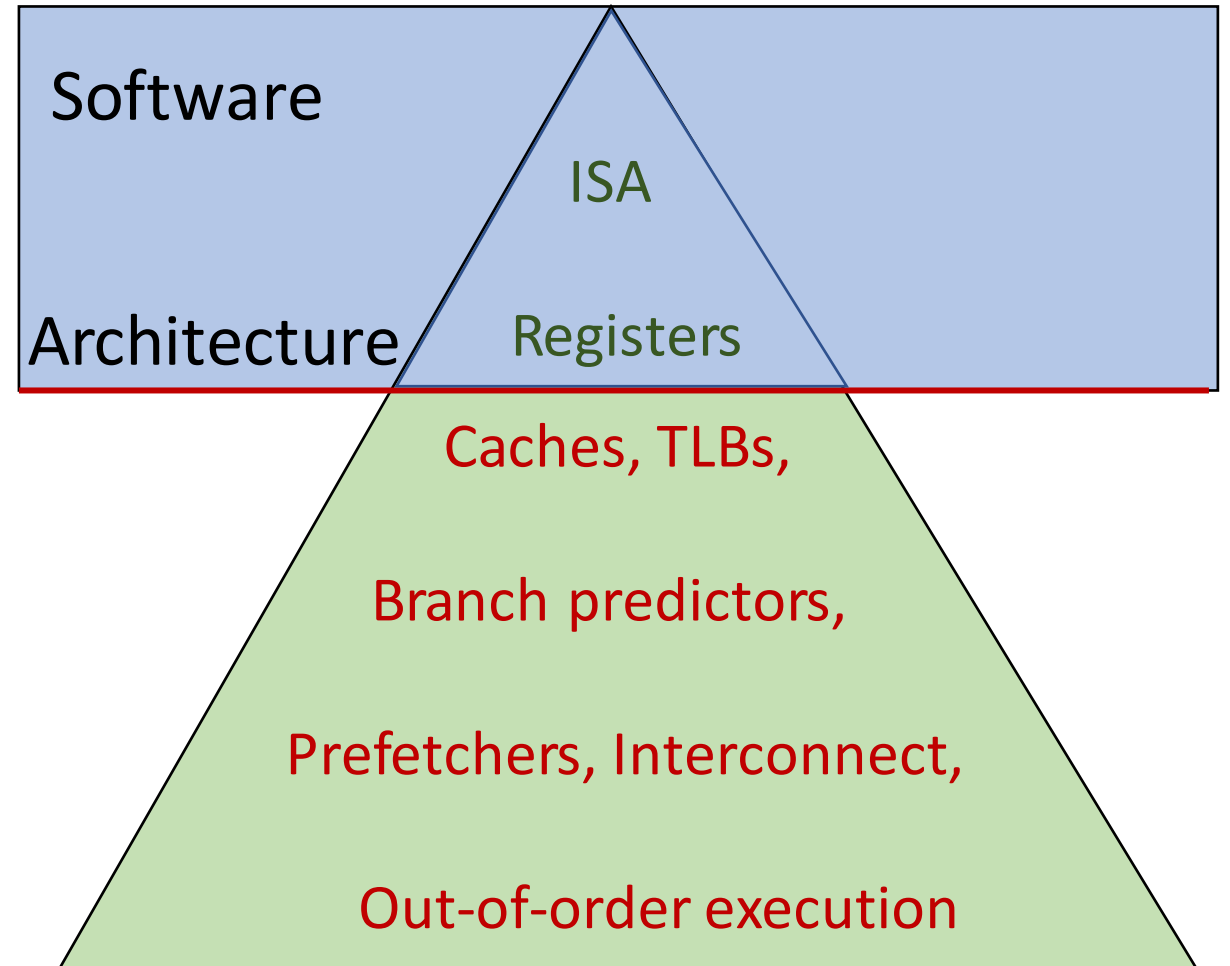


Not exposed to programmer





Memory

From Performance to Security: 10K Feet View





CAWS Talks and Ideas: Secure?




CAWS Talks and Ideas: Secure?

Attacker




Attacker



O3 processor ka, cache ka, prefetcher ka, coherence protocol ka, DRAM/NVRAM ka, sabka khabar lega 😞

What Is Security?



CIA ☹️

Security: A bit Subtle

Confidentiality

*You do not **see (READ)** what you are not supposed to see*

Integrity

*You do not **change (WRITE)** what you are not supposed to see*

Availability

*You do not **affect (DELAY)** others (un)intentionally*

Attacks Inside



inside™




inside™



inside™

Media Articles ☹️

New SWAPGS Side-Channel Attack Bypasses Spectre and Meltdown Defenses



Hackers Can Mess With Voltages to Steal Intel Chips' Secrets

A new attack called Plundervolt gives attackers access to the sensitive data stored in a processor's secure enclave.




'RAMBleed' Rowhammer attack can now steal data, not just alter it

Academics detail new Rowhammer attack named RAMBleed.



By Catalin Cimpanu for Zero Day | June 11, 2019 -- 17:00 GMT (22:30 IST) | Topic: Security



Intel Responds to ZombieLoad and CacheOut Attacks


By Nathaniel Mott 4 days ago

not just with a "what?"



The Elite Intel Team Still Fighting Meltdown and Spectre

One year after a pair of devastating processor vulnerabilities were first disclosed, Intel's still dealing with the fallout.



NEW SIDE-CHANNEL
ATTACK EXTRACTS
PRIVATE KEYS FROM
SOME QUALCOMM
CHIPS

Apr 18, 2019

By Dennis Fisher

Brushing-up: Information Leakage

$x \leftarrow 1$

Modular exponentiation, $b^e \bmod n$

for $i \leftarrow |e|-1$ **downto** 0 **do**

Exponent e is used for decryption

$x \leftarrow x^2 \bmod n$

square

if ($e_i = 1$) **then**

reduce

$x = xb \bmod n$

endif

multiply


done

return x

$e_i = 0$, Square Reduce (SR)
 $e_i = 1$, SRMR

Attacker tries to get the e

Timing Channel



Adpated from flush+reload attack [Usenix Security '14]


Toy Example: Flush Based Attacks

```
If secret=1 do  
    access(&a)  
else // secret=0  
    no-access
```

Victim

```
flush(&a)  
t1=start_timer  
    access(&a)  
t2=end_timer
```

Attacker




Fast – 1

Slow – 0

Side and Covert Channels

Spy




Side-channel attacks

Victim



Let's
play




Covert-channel attacks

Oh Yes!!



Shared LLC Attacks

Attacks at the LLC exploit timing channels:
LLC miss > LLC hit



Flush + Reload

Evict + Reload

Prime + Probe

clflush


Eviction based attacks

Threat Model



Knowing the victim *has accessed a cache set (line)* can be considered as a *successful* attack

Flush+Reload Attack



Step 0: *Spy maps the shared library, shared in the cache*



Flush+Reload Attack



Cflush




LLC

Step 0: *Spy maps* the shared library, shared in the cache

Step 1: *Spy flushes* the cache block



Flush+Reload Attack




Step 0: *Spy maps* the shared library, shared in the cache

Step 1: *Spy flushes* the cache block

Step 2: *Victim reloads* the cache block



Flush+Reload Attack



Step 0: *Spy maps* the shared library, shared in the cache


Step 1: *Spy flushes* the cache block

Step 2: *Victim reloads* the cache block

Step 3: *Spy reloads* the cache block (hit/miss)



Flush + Flush




Step 0: *Spy maps the shared library, shared in the cache*

Step 1: *Spy flushes the cache block*



Flush + Flush




Step 0: Spy *maps* the shared library, shared in the cache

Step 1: Spy *flushes* the cache block

Step 2: Victim *reloads* the cache block



Flush + Flush



Step 0: Spy *maps* the shared library, shared in the cache

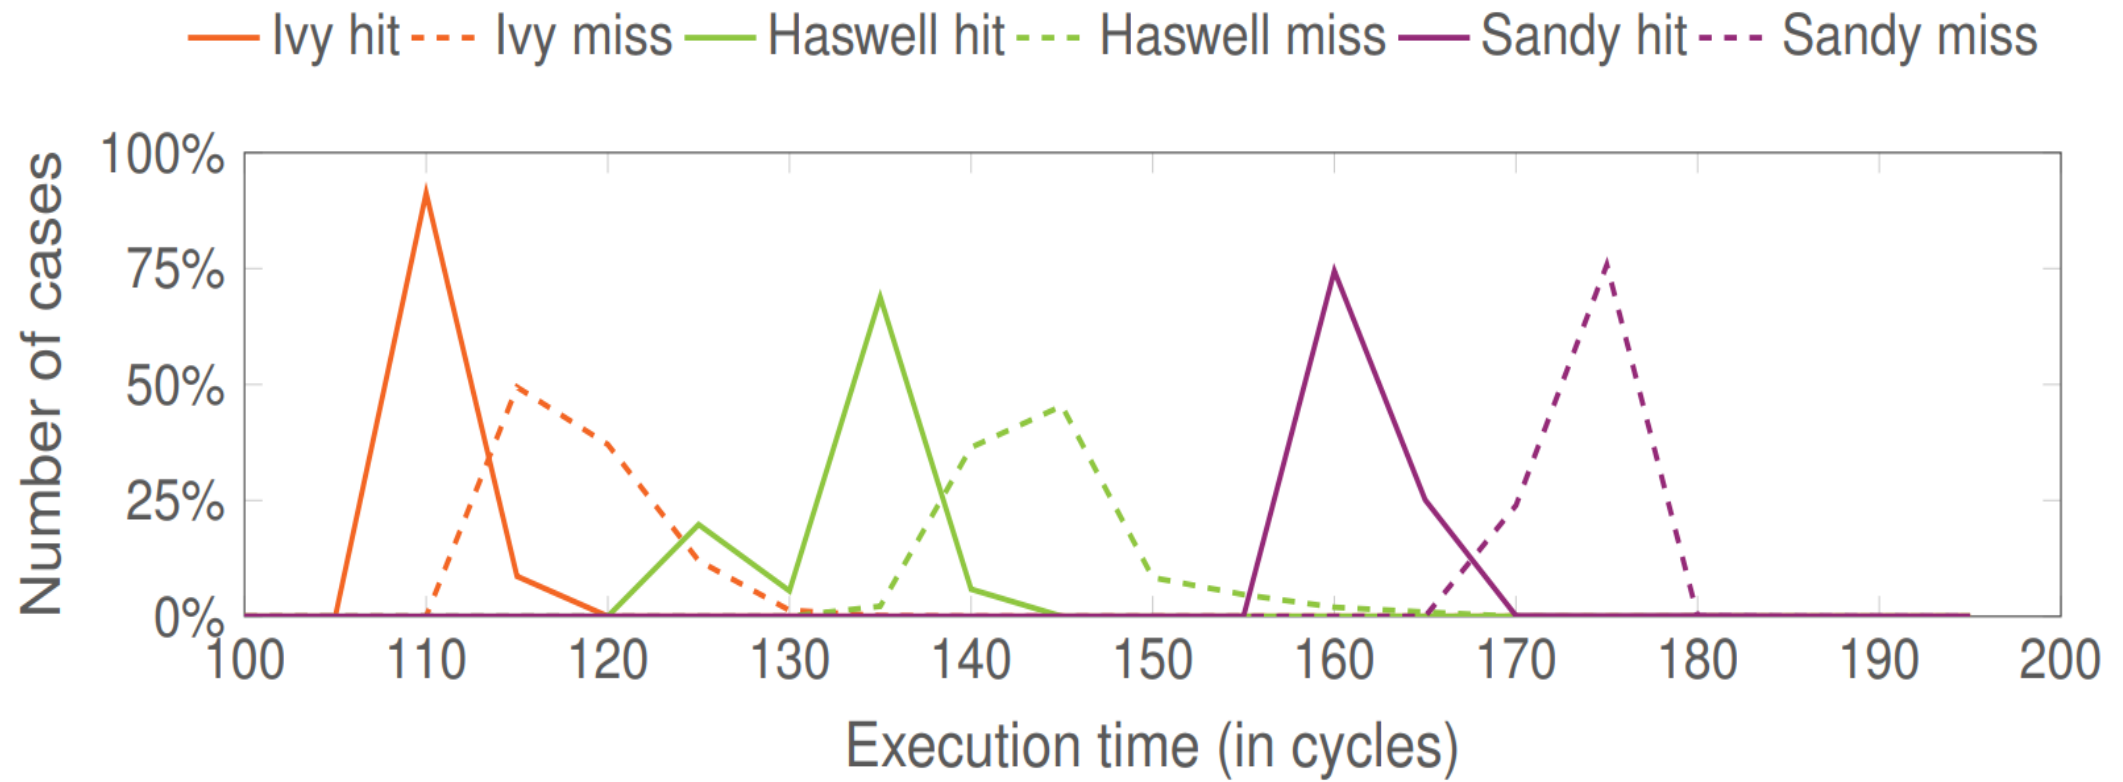
Step 1: Spy *flushes* the cache block

Step 2: Victim *reloads* the cache block

Step 3: Spy *flushes* the cache block again



Confused?



No sharing?

What If I do not share anything with you ??



Do not worry, I have Amazon Prime

amazon.com
Prime

Whaaaaat?!



Sorry:


amazon.com
Prime+Probe



Prime+Probe



Step 0: *Spy fills the entire shared cache*



LLC




Prime+Probe



Step 0: Spy *fills* the entire shared cache


Step 1: Victim *evicts* cache blocks while running



LLC



Prime+Probe



Step 0: Spy *fills* the entire shared cache

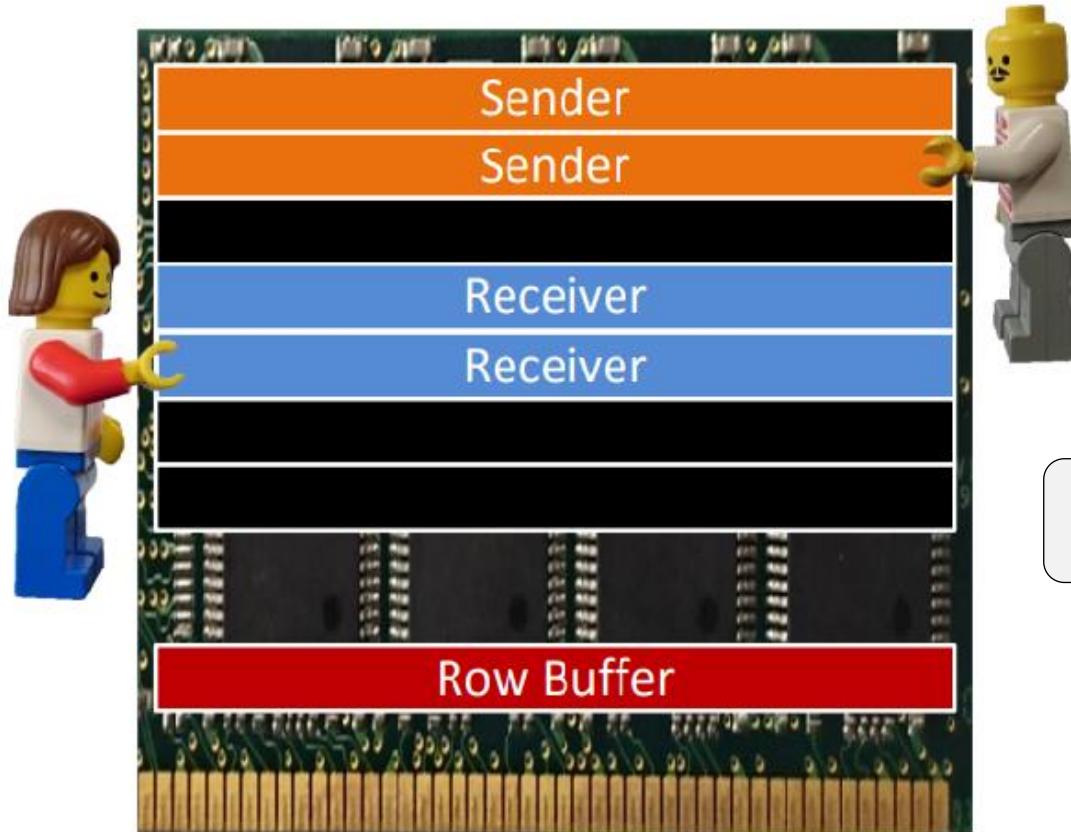
Step 1: Victim *evicts* cache blocks while running

Step 2: Spy *probes* the cache set

If misses then victim has accessed the set



Same at the DRAM




Row buffer hit/miss

Same at all the shared resources

Is it that Simple?


Notion of Time Gap




Inclusiveness




L1/L2



LLC




Miss



Inclusiveness




L1/L2



Cross-core back-invalidation

LLC




Miss

Inclusiveness




Attacker knows whether victim has accessed a set or not

L1/L2



Miss

LLC



Miss

System Noise



DVFS aware

Aware of co-location of victim

Fast and Stealthier

High accuracy even on noisy systems



<https://car3s.github.io/dabangg/>

The Hacker News

**New Noise-Resilient Attack On Intel and AMD CPUs
Makes Flush-based Attacks Effective**



New Technique Improves Effectiveness of Timing Channel Attacks

By [Ionut Arghire](#) on June 01, 2020



Share



Tweet



Recommend 0



RSS

Two researchers have discovered a new timing channel attack technique that remains effective even if multiple processes are running on a system.


Called DABANGG (the Hindi word for fearless), the newly proposed technique improves the effectiveness of flush-based attacks such as Flush+Reload and Flush+Flush, researchers Anish Saxena and Biswabandan Panda from the Indian Institute of Technology Kanpur claim in a [research paper](#).

How Practical?



Future is uncertain, if we do not take care of present attacks, future may be worse ☹️

Spectre and Meltdown




Spectre in Action: Fasten Your Seat Belts

```
int CAWSArray = [100, 200, 300];  
int attacker = 4;  
if (attacker < sizeof(CAWSArray))  
    y = MyArray[CAWSArray[attacker]*512]
```

DRAM LOAD

DRAM LOAD



Branch Predictor and Speculative Execution

```
int CAWSArray = [100, 200, 300];  
int attacker = 4;  
if (attacker < sizeof(CAWSArray))  
    y = MyArray[CAWSArray[attacker]*512]
```



Branch predictor returns TRUE ☹️

Branch Predictor and Speculative Execution

```
int CAWSArray = [100, 200, 300];  
int attacker = 4;  
if (attacker < sizeof(CAWSArray))  
    y = MyArray[CAWSArray[attacker]*512]
```



T T T T T T T T T T

Branch predictor returns TRUE ☹️

Attacker has mis-trained it ☹️ ☹️

How? By using values less than 3 always ☹️ ☹️

Branch Predictor and Speculative Execution

```
int CAWSArray = [100, 200, 300];  
int attacker = 4;  
if (attacker < sizeof(CAWSArray))  
    y = MyArray[CAWSArray[attacker]*512]
```

Branch predictor returns TRUE ☹️

Attacker has mis-trained it ☹️ ☹️

Processor is on the wrong-path ☹️ ☹️ ☹️

Branch Predictor and Speculative Execution

```
int CAWSArray = [100, 200, 300];  
int attacker = 4;  
if (attacker < sizeof(CAWSArray))  
    y = MyArray[CAWSArray[attacker]*512]
```

Branch predictor returns TRUE ☹️

Attacker has mis-trained it ☹️ ☹️

Processor is on the wrong-path ☹️ ☹️ ☹️

Branch resolution latency 200 cycles ☹️ ☹️ ☹️ ☹️

Within these 200 cycles 😊

```
int CAWSArray = [100, 200, 300];  
int attacker = 4;  
if (attacker < sizeof(CAWSArray))  
    y = MyArray[CAWSArray[attacker]*512]
```

CAWSArray[4] is in L1/L2/L3 ☹️

The address is in the cache ☹️ ☹️

Yes, you guessed it right: F+R, P+P cache attacks ☹️ ☹️ ☹️

Picture Abhi Baki Hai 😊 After 200 cycles

Processor realized it was a mistake and *flushed* all wrong path instructions

But cache has the data 😞

*y = MyArray[CAWSArray[attacker]*512]*

LOAD MyArray[0] 60 ns

LOAD MyArray[512] 60 ns

LOAD MyArray[1024] 5 ns Bingo !! CAWSArray[attacker] = 2



Meltdown: The O3 Curse!!

```
1. raise_exception();  
2. // line below is never reached  
3. secret=KernelArray[data*4096];
```

Kernel Trap

```
1. secret=KernelArray[data*4096];  
2. raise_exception();
```

Out-of-order (O3) as
it has no dependency

What about page-fault?

Mitigations: Read it On your Own !!

Attacks in the Cloud



SIDE-CHANNEL AND **COVERT**
CHANNEL ATTACKS AT THE **CACHE**
AND THE **DRAM IN CLOUD**




Attacks in the Cloud



I use cloud

No fear of information leakage ??









server



CPU #1



CPU #2



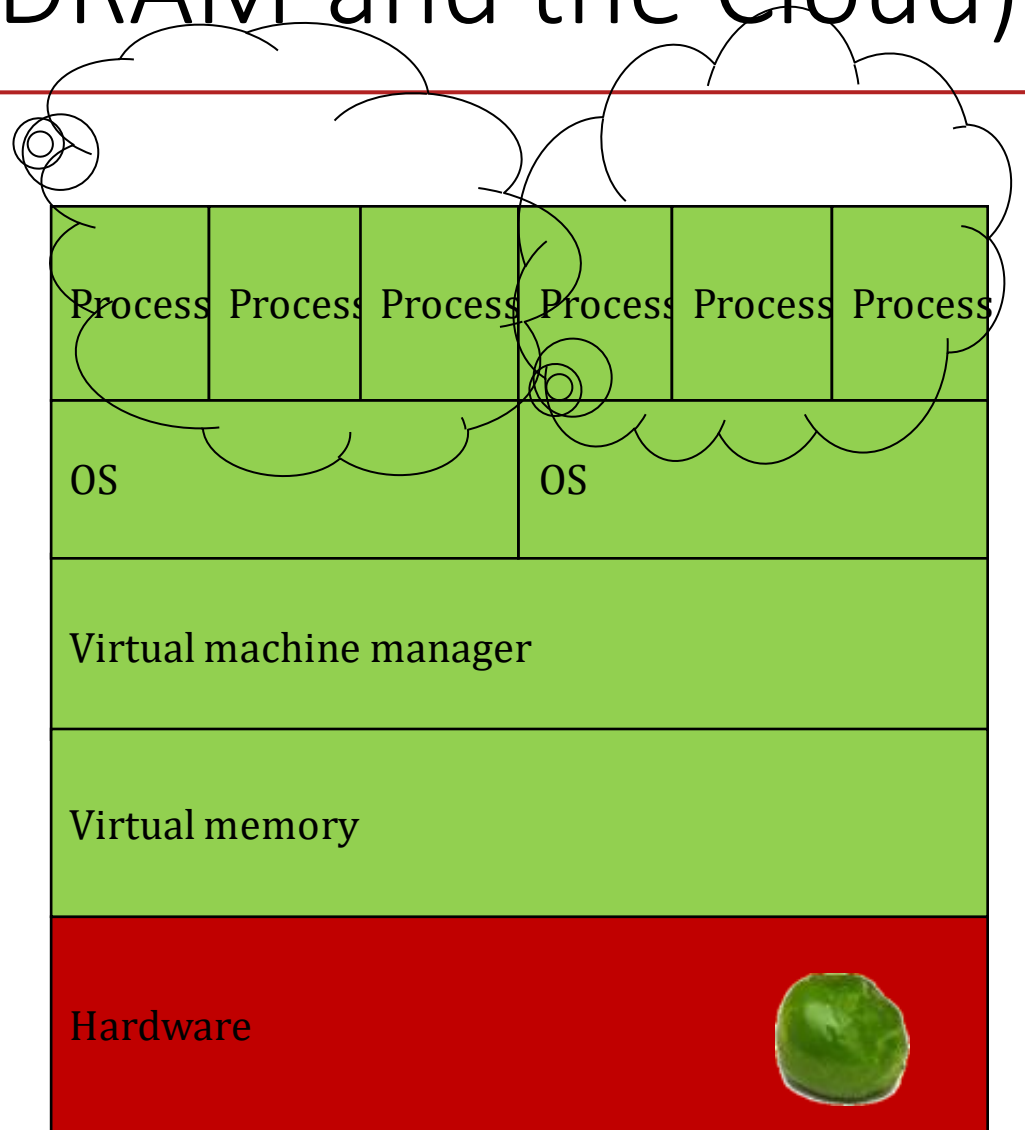

DRAM



Gotcha!!

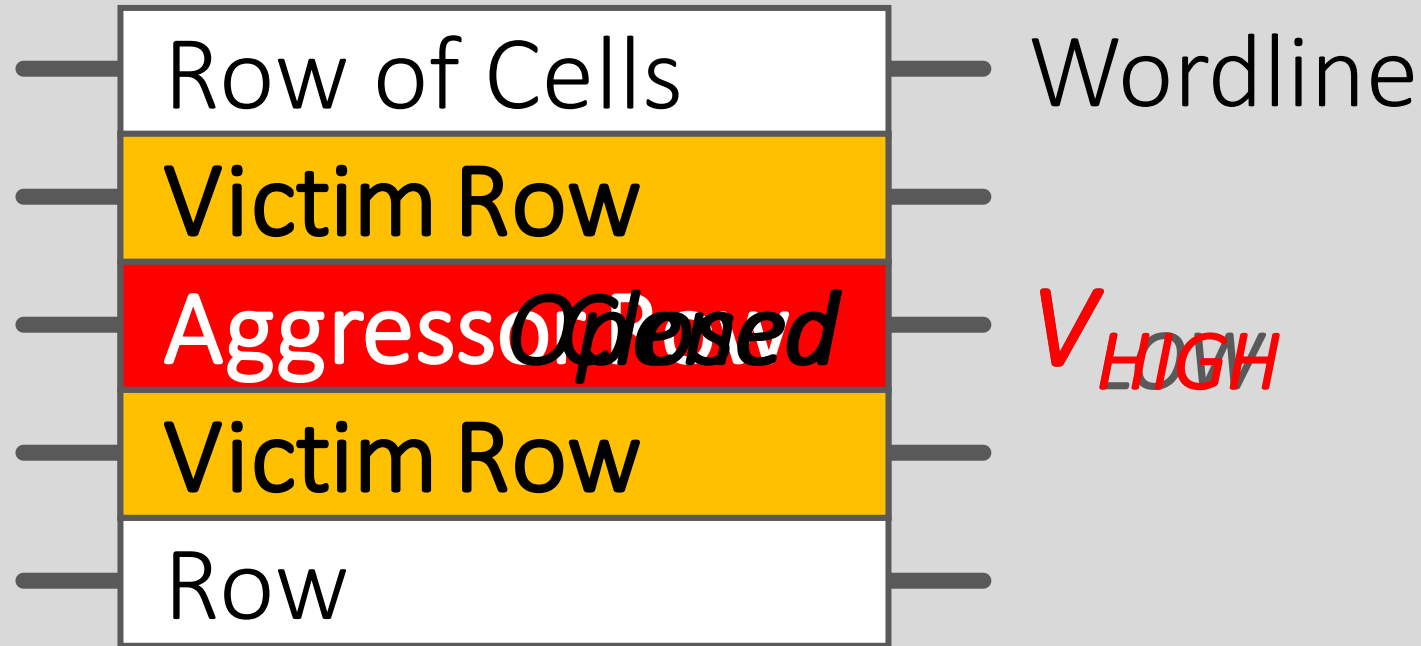
[Courtesy: Eran Tromer]

Pea and the Princess (LLC/DRAM and the Cloud)

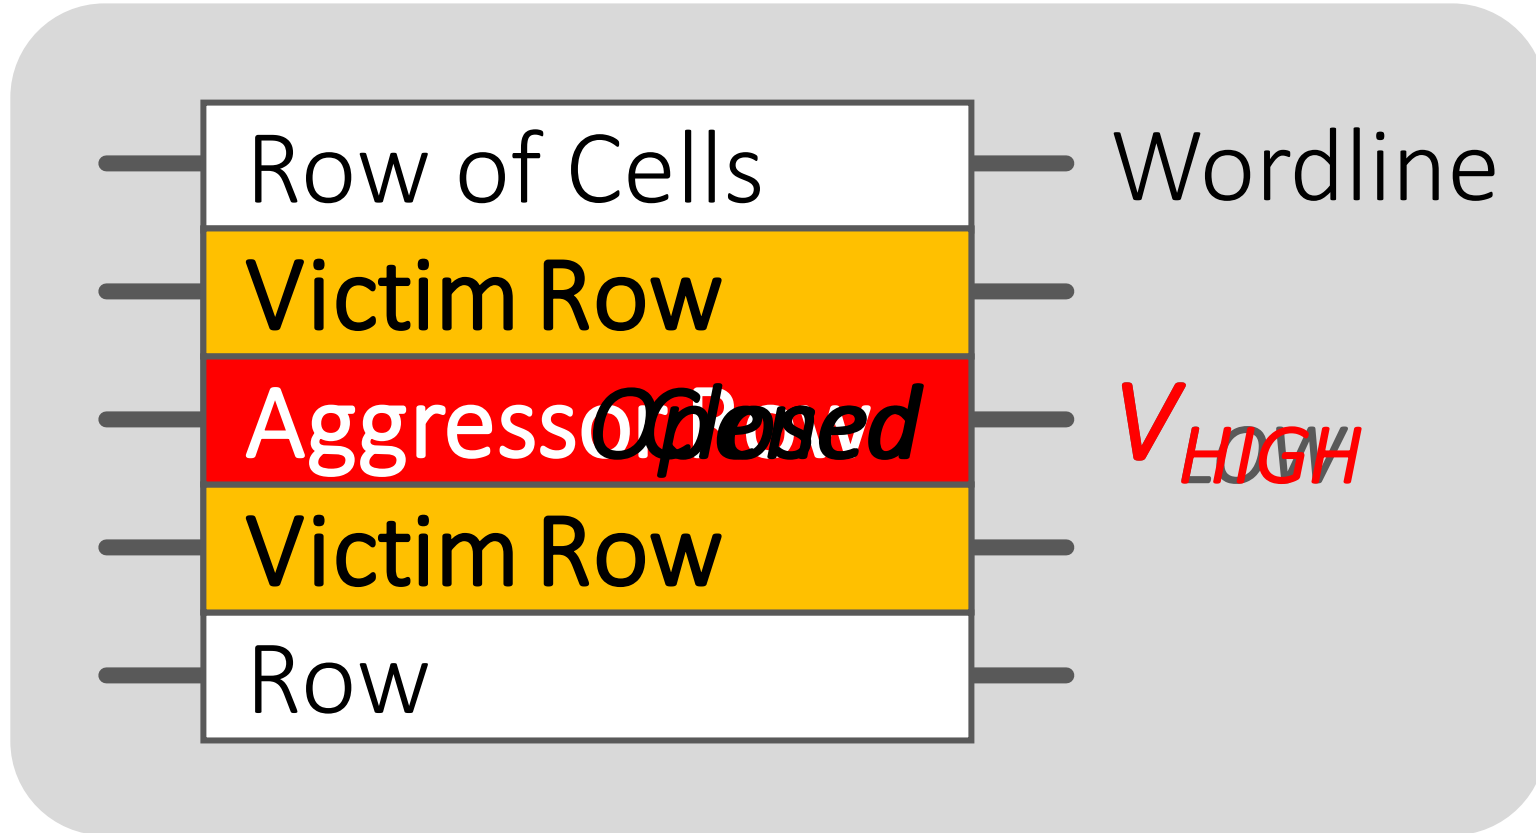


RowHammer (Time for Integrity based attacks)

Adpated from Row Hammer attack [ISCA '14]




*Repeatedly opening and closing a row induces **disturbance errors** in adjacent rows*




“It’s like breaking into an apartment by repeatedly slamming a neighbor’s door until the vibrations open the door you were after” – Motherboard Vice

RowHammer




1. Avoid *cache hits*
 - Flush **X** from cache
2. Avoid *row hits* to **X**
 - Read **Y** in another row



RowHammer (But Why? Read it Later)

```
loop:  
  mov  (X),  %eax  
  mov  (Y),  %ebx  
  clflush (X)  
  clflush (Y)  
  mfence  
  jmp  loop
```



Some Cool Attacks




Before powering off

Freeze it to -50°C

Cool It Down (Data remanence, cold-boot attack)




Put It Back



After Five Seconds



After 30 Seconds




After 300 Seconds



“Think about non-volatile memory”



DOS Attacks: At the DRAM (Availability)



How?

```
// initialize large arrays A,  
B
```

streaming

```
for (j=0; j<N; j++) {  
    index = j*linesize;  
    A[index] = B[index];  
    ...  
}
```

STREAM

Sequential memory access

high row buffer locality (96% hit rate)

Memory Intensive

```
// initialize large arrays A,  
B
```

random

```
for (j=0; j<N; j++) {  
    index = rand();  
    A[index] = B[index];  
    ...  
}
```

RANDOM

Random

Low row buffer locality (3% hit rate)

Memory Intensive

Trusted Execution Environments

arm
TRUSTZONE



Not **resistant** to timing channels

Sky is Certainly Not Falling



Media articles *sensationalize* ☹️

All the attacks are POCs

Can we make these attacks real?

Offensive side



Microarchitecture need to be fixed to avoid surprises

Can we design secure microarchitecture?

Defensive side

Research Questions

- Relatively new field of research.
- Top forums: USENIX SECURITY, S&P, CCS, NDSS, WOOT along with top computer architecture conferences.
- Performance-security trade-offs
- New attacks, new threat models
- Cat-mouse game 😊

Secure Thanks



Ack

- Some of the slides are adapted and modified from Clementine Maurice