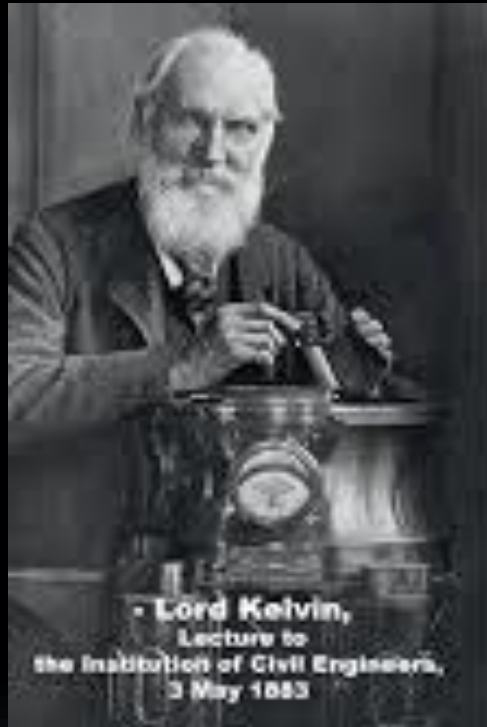


K V Subramaniam
Center for Cloud Computing and Big Data
PES University, Bangalore
subramaniamkv@pes.edu



PERFORMANCE TOOLS

COMPUTER ARCHITECTURE WINTER SCHOOL - 2020



**"I often say that
when you can measure
what you are speaking about,
and express it in numbers,
you know something about it;
but when you cannot measure it,
when you cannot express it in numbers,
your knowledge is of a meagre
and unsatisfactory kind;
it may be the beginning of knowledge,
but you have scarcely in your thoughts
advanced to the state of Science,
whatever the matter may be."**





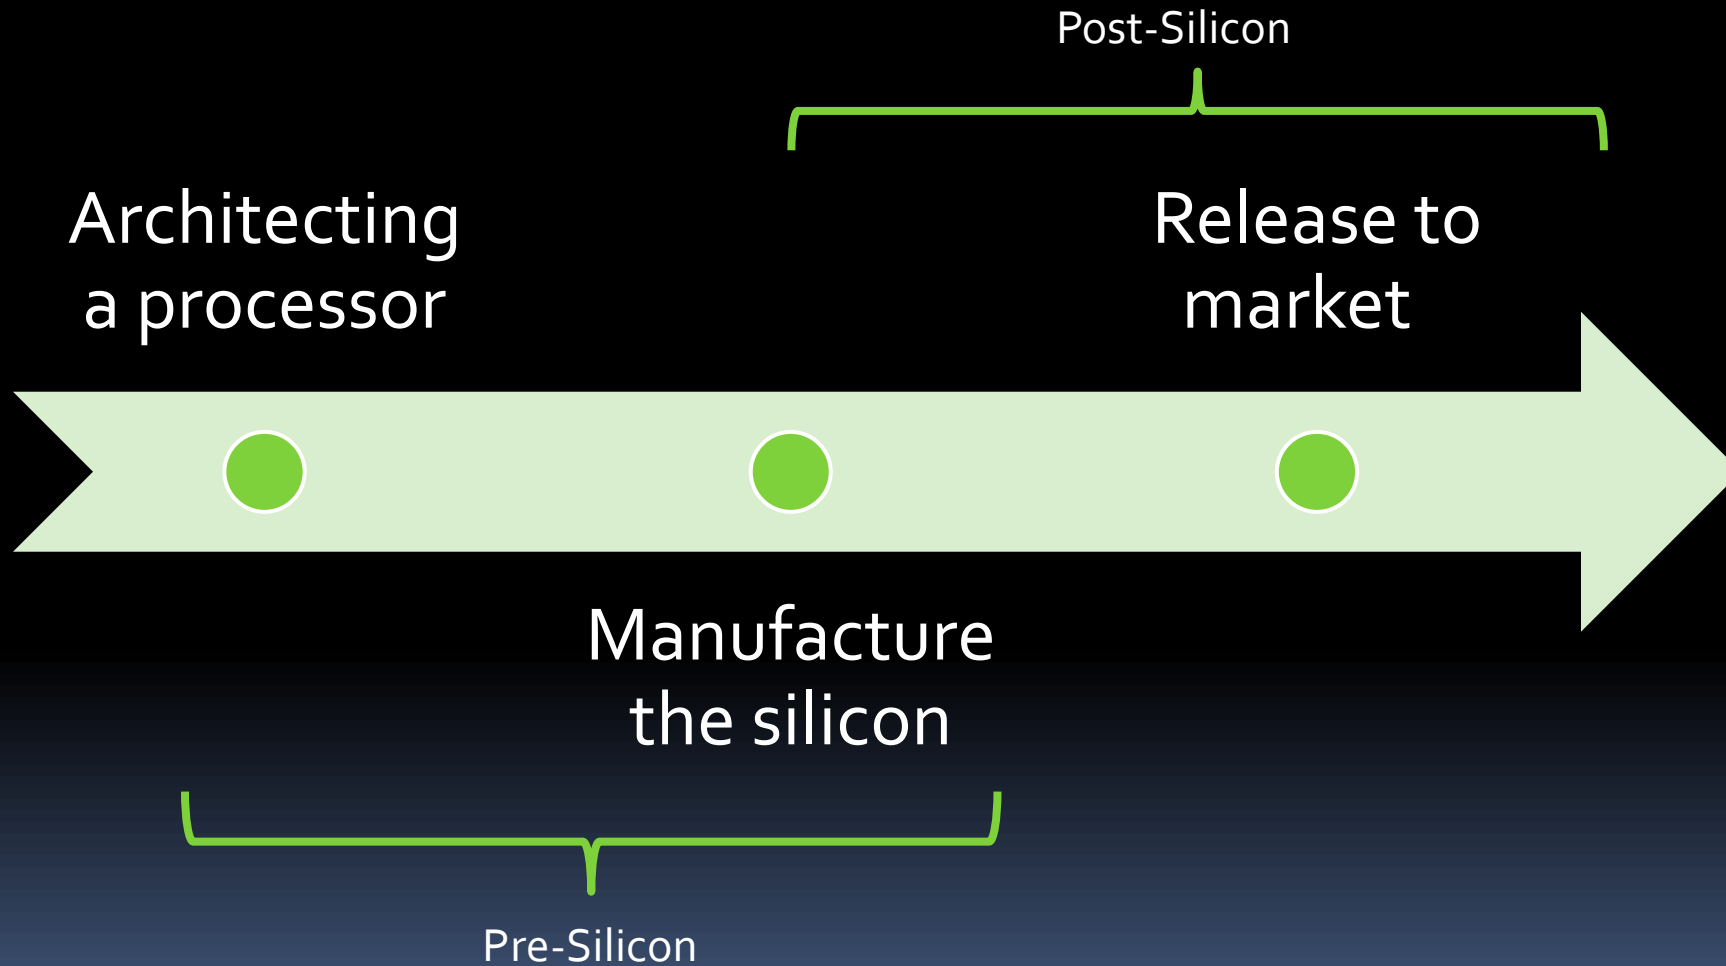
Overview



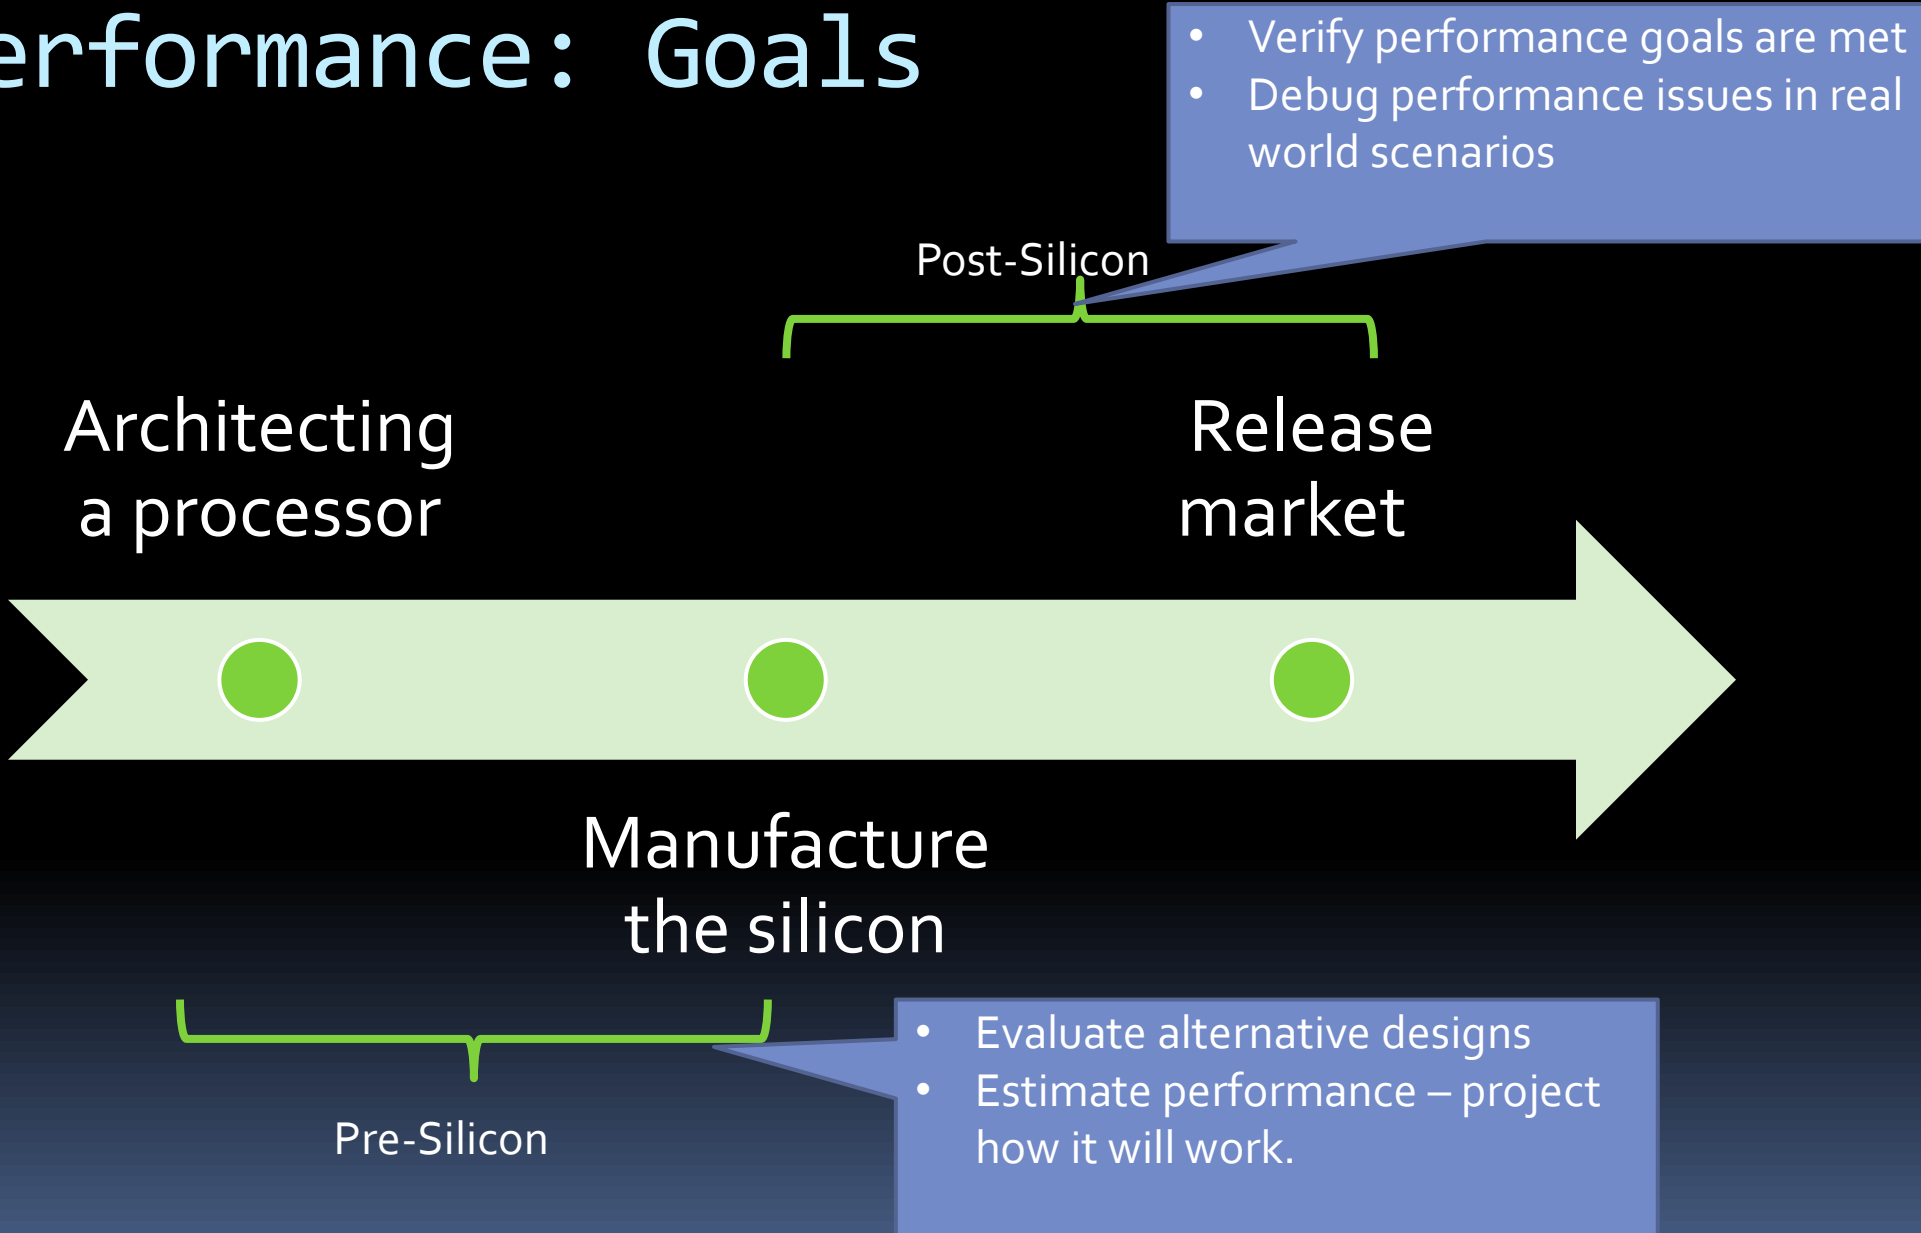
- Why performance?
- Performance metrics
- Simulators
- Profiling programs – CPU/Memory
- Measuring Operating system activity
- Measuring microarchitectural activity



Background



Performance: Goals





Overview



- Why performance?
- Performance metrics
- Simulators
- Profiling programs – CPU/Memory
- Measuring Operating system activity
- Measuring microarchitectural activity



Performance Metrics

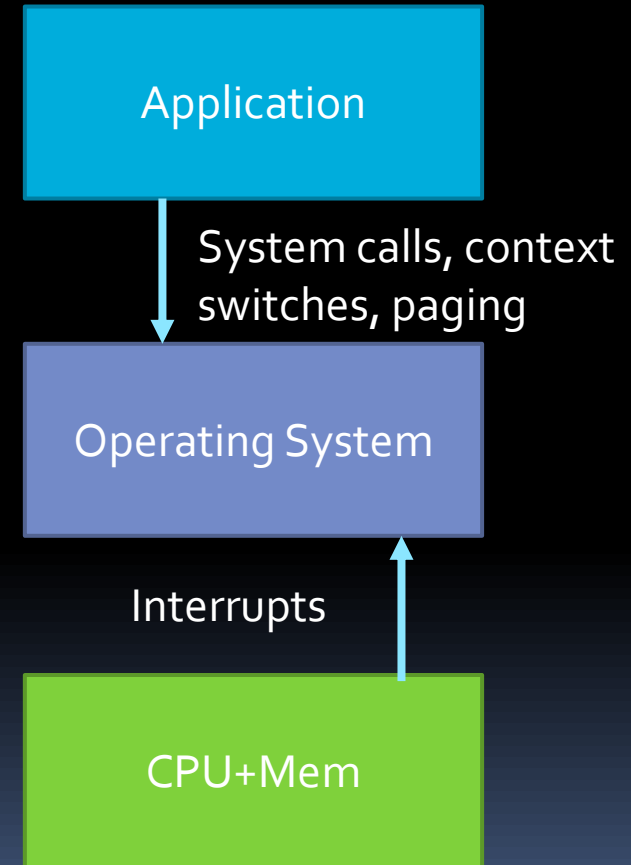
- How fast is the application running?
 - #clock cycles
 - Is it using the architecture efficiently?
 - Instructions per Cycle (IPC)
- Pre-silicon – how to check?
- What factors does application performance depend on?
- How do we know that it is using the architecture efficiently?

Application performance factors

- Algorithm used
 - From a theoretical standpoint this is important.
- For us more important
 - Where is time being spent?
 - How does it map to the hardware?

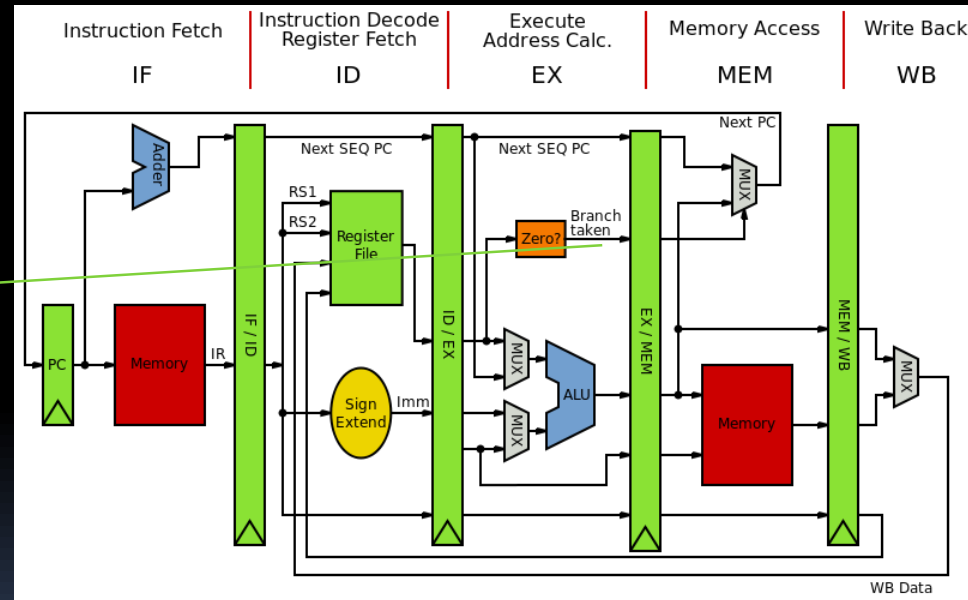
Software stack impact

- Compiler – maps applications to the hardware – multi-dimension arrays
- Runtime systems – OS, libraries, JVM, Garbage collectors
- Applications interact with OS
 - Voluntarily – when they need resources
 - Involuntarily – due to interrupts/context switches.
- Performance can depend on OS activity
 - Other applications..
 - Paging/TLB structures
- Becomes more complex with addition of
 - Virtual machines, containers.
 - And newer programming models – FaaS, Microservices.



Performance Metrics .. 2

- For a computer architect, knowing IPC itself is not enough
 - We need a deeper insight into how the various components are performing
 - e.g – Branch predictors, caches, prefetchers etc.
 - Ferdman et. al ASPLOS 12 – compared scale out workloads with SPEC





Overview

- Why performance?
- Performance metrics
- **Simulators**
- Profiling programs – CPU/Memory
- Measuring Operating system activity
- Measuring microarchitectural activity

Simulators - Pre-silicon

- Model the system being analyzed
 - Pipeline(s), caches, cores...
 - Both individually and combined together
- Use of simulators to
 - Check alternative designs
 - Interactions between various components
 - Performance projections on standard workloads

Types of simulators

- Functional Simulators

- Check if the hardware design is indeed working correctly
- ISA simulators
- Check if the software stack will run correctly.

Add r1, r2[1000]

Adds content of
memory to register r1

Needs to keep track of
the **state** of the system

Types of simulators – Timing Performance simulators

- Timing/Performance simulators
 - How long does an (set of) instruction(s) take to execute?
 - Not so much interested in the results per se.
 - Accuracy – cycle accurate simulators

Add r1, r2[1000]

Compute total cycles..

- Check if location is DRAM/cache
- Move data to cache and then to register
- Perform the add
- Also keep track of how much time is spent in each activity



Overview

- Why performance?
- Performance metrics
- Simulators
- Profiling programs – CPU/Memory
- Measuring Operating system activity
- Measuring microarchitectural activity



Profiling programs

```
Func Add()  
{  
  
}  
  
Func Multiply()  
{  
  
}  
  
Func Main()  
{  
    call Add()  
    call Multiply()  
}
```

- Where is the time being spent?
- Techniques
 - ▣ Instrumentation
 - ▣ Sampling

Instrumentation based profilers

```
Func Add()  
{  
  
}  
  
Func Multiply()  
{  
  
}  
  
Func Main()  
{  
    call Add()  
    call Multiply()  
}
```

Add additional code
To each function

Log – entry exit points

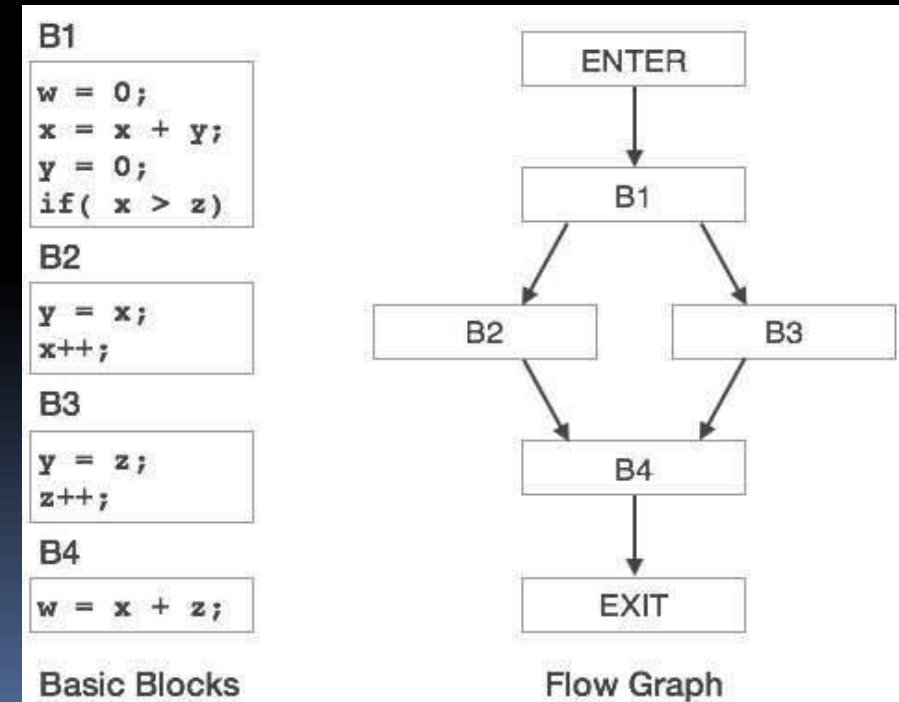
```
Func Add()  
{  
    print(...)  
}  
  
Func Multiply()  
{  
    print(...)  
}  
  
Func Main()  
{  
    print(...)  
    call Add()  
    call Multiply()  
}
```

Instrumentation based profilers

- Granularity of instrumentation
 - Function level – prof/gprof
 - High level identification of functions that take time
 - Basic block level
 - Unit of code with a single entry and exit
 - Additional instrumentation added to each basic block
 - detailed instruction execution counts.
 - Code paths

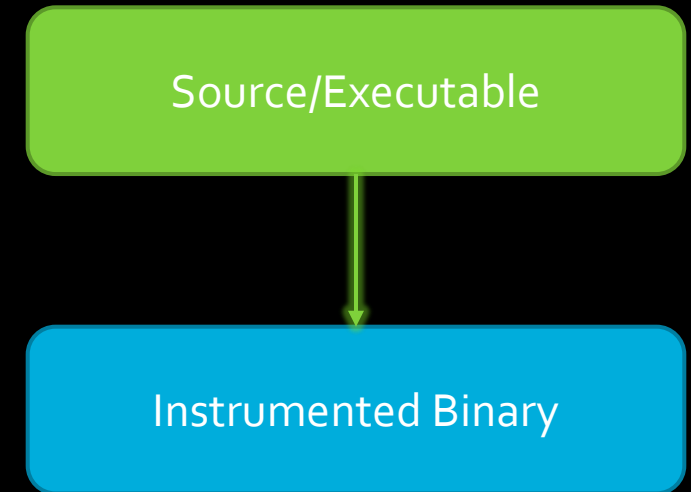
Image source:

https://www.tutorialspoint.com/compiler_design/compiler_design_code_optimization.htm



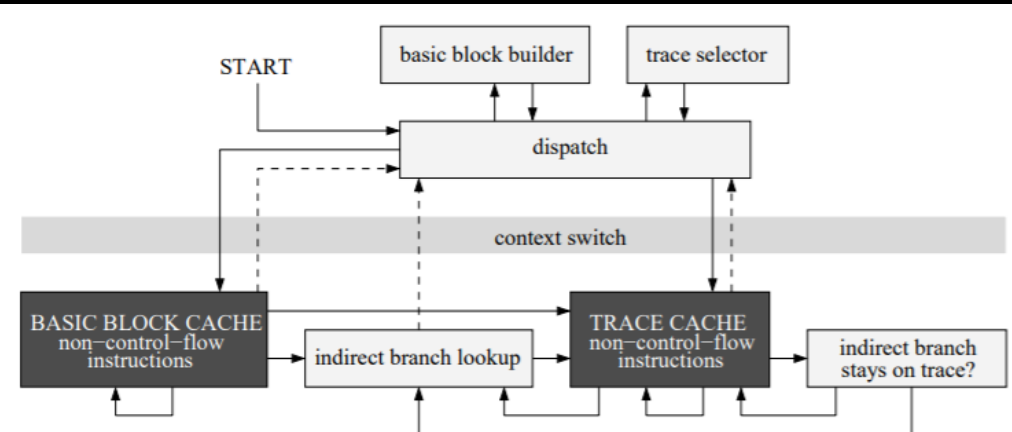
When to instrument?

- Static
 - compile time or prior to execution
 - Gprof: `cc helloworld.c -pg`
 - Produces an instrumented a.out.
 - Executing this produces run time profile data
 - Analyzed by post processing tools
 - Call graphs



When to instrument?

- Dynamic – at run time.
 - Examine instructions/basic blocks before they run
 - Add instrumentation code and translate to new address space.
 - Execute the translated binary
 - Cache translations
 - E.g: DynamoRIO
 - Valgrind
 - Memory leaks, data races

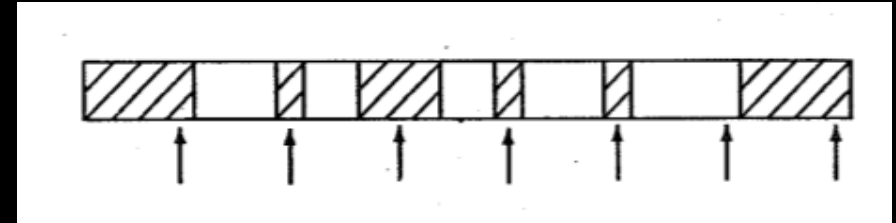


Caveats: instrumentation based

- Addition of instrumentation code
 - Code bloat
 - Additional code executed dynamically
 - Not really good for measuring efficiency of hardware usage – cache misses?
- Dynamic instrumentation – heavy overheads during run time
 - Good for generating detailed instruction/data traces
 - These can be useful for input to simulators

Sampling based profilers

- Interrupt program at regular intervals
- Examine call stack
- Identify the function currently being executed.
- On completion,
 - $\text{Fraction of samples} / \text{total \#samples} \rightarrow$ approximated as runtime spent in function.



Sampling based profilers..

- Sampling frequency affects accuracy.
- Short lived functions may not show up in sample
- OS activity – interrupts can affect accuracy
- Good for estimating overall breakup of time
 - Fast, less intrusive
- Not so good if we are looking for detailed breakup or causes of latency.



Overview

- Why performance?
- Performance metrics
- Simulators
- Profiling programs – CPU/Memory
- Measuring Operating system activity
- Measuring microarchitectural activity

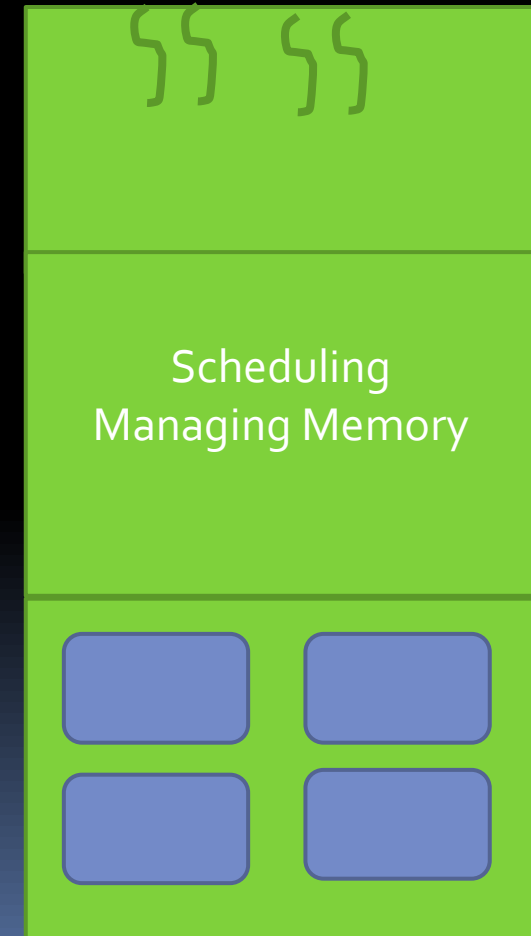
Where does OS impact

Applications are multi threaded
Have large memory footprint

OS – Manages mapping, resources

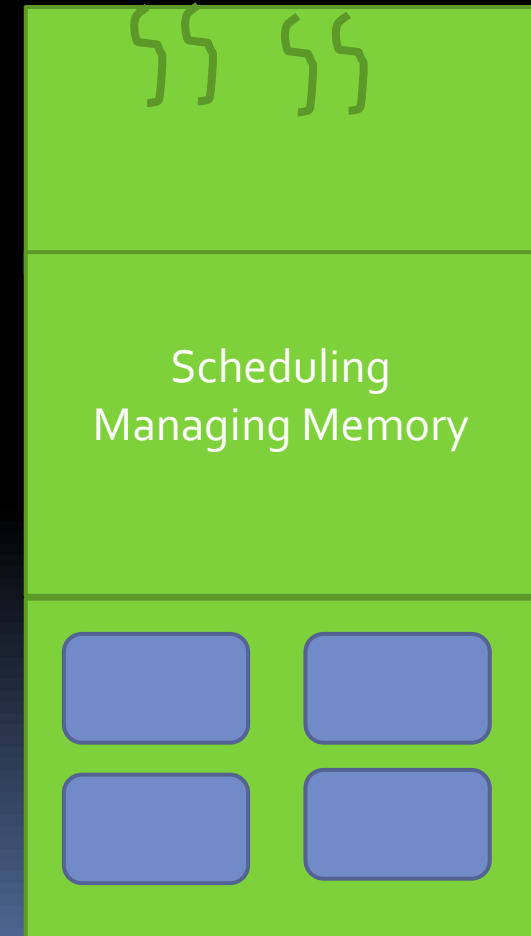
Can lead to inefficient mapping

Architecture: Multicore, shared
memory/cache
Resource sharing



Thread Mapping

- A thread executes on a core, it builds up context
 - Cache : enables fast access to data/instruction
 - Memory: pages brought into memory
- On context switch, thread can migrate to different core
 - Context is lost – memory context is important – NUMA machines
- Thread – Core mapping
 - htop
 - Shows the current thread being executed on a core.
- How to measure other activity of a process?



Peeking into the OS

- /proc filesystem
 - A virtual filesystem
 - Gives access to the kernel data structures
 - e.g Virtual memory mapping, open files,
- Sar – system activity report
 - Measure paging activity
 - Disk activity
 - Network activity...




Overview

- Why performance?
- Performance metrics
- Simulators
- Profiling programs – CPU/Memory
- Measuring Operating system activity
- Measuring microarchitectural activity





Measuring Microarchitectural activity

- Sometimes we really need to understand
 - How well application uses hardware.
 - For example:
 - Caches, TLBs,
 - Need accurate information about such events
- 

Hardware counters

- Hardware maintains counters of events
 - Introduced in Pentium
(<http://archive.gamedev.net/archive/reference/articles/article213.html>)
- Many events within the hardware are measured
 - Available as MSR (Model Specific Registers)
 - What can be measured depends on the model
 - RDMSR/WRMSR to read/write registers
- Tools: Intel Vtune, perf, AMD uPerf

Using Performance Counters

- Threads from two different cores sharing L3
 - Can interfere with each other
 - Remove cache lines required by the other.
- Perf can be used to identify the number of accesses made by each core.

Instruction Based Sampling

- Given modern processors perform out of order execution
- Many instructions will be simultaneously be active
 - In different stages of execution.
 - Just measuring stats does not allow attributing it to a specific instruction.
- IBS(AMD) – allows sampling certain instructions
 - Collecting all stats about them – during fetch and execution phases.
 - For example: how many times was a speculative fetch for an instruction killed

Conclusion

- Performance
 - Important in many parts of architecture design
 - Needs knowledge of different aspects of system – compilers, OS ...

- Reminder – Lab manual will be shared on Channel